

Programming Assignment #2: Linear Models & Optimization

Due date: October 28th, 2016 (Friday) 4:10PM before class

“On my honor, as an Aggie, I have neither given nor received unauthorized aid on this academic work.”

Signature:

Name:

0. Reading Assignments: The relevant content in BRML Chapters 17 & Appendix

1. Implement the gradient descent/ascent algorithm for logistic regression (40pts):

1. Use the data set with binary classes (**bclass-train** and **bclass-test**): Each row is one example with the first column as the class label $\{-1, 1\}$ and the rest as feature variables) You can replace the class label -1 by 0 . You can find the data set in the course content.
2. Run your logistic regression using the raw data, data that has been normalized to have unit l_1 norm, and data that has been normalized to have unit l_2 norm. Which seems to work best?
3. Plot error rates for the training, and test data as a function of iteration (both the raw predictions and the normalized predictions).

2. Locally Weighted Logistic Regression (40pts):

Implement a locally-weighted version of logistic regression, where we weight different training examples differently according to the query point. The locally weighted logistic regression problem is to maximize

$$l(\beta) = \sum_{i=1}^N w^i \left\{ y^i \log f_{\beta}(x^i) + (1 - y^i) \log [1 - f_{\beta}(x^i)] \right\} - \lambda \beta^T \beta,$$

where the last term is the regularization term as we discussed for the linear regression in class. (You can also implement the regularized logistic regression for **1**, which often can give more stable results using either gradient descent or Newton’s method.) You can set $\lambda = 0.001$; Or you can use the development data included in the data set to pick the best performing λ :

1. Compute the gradient $\nabla_{\beta} l(\beta)$ and the Hessian matrix $H = \nabla_{\beta} [\nabla_{\beta} l(\beta)]$;
2. Given a new test data point x , we compute the weight by

$$w^i = \exp\left(-\frac{\|x - x^i\|_{l_2}^2}{2\tau^2}\right),$$

where τ is the bandwidth. Use the same data set with binary classes as above (**bclass-train** and **bclass-test**) to implement **Newton’s** method for this locally weighted logistic regression. Vary $\tau = \{0.01, 0.05, 0.1, 0.5, 1.0, 5.0\}$ to: (a) compute w^i ’s for each development/test sample using the formula above, (b) maximize $l(\beta)$ to learn β , (c) predict y based on $f_{\beta}(x)$ ($y = 1$ when $f_{\beta}(x) \geq 0.5$), and finally (d) plot the error rates with respect to τ and compare them with the ones obtained in **1**.

3. Support Vector Machines (20pts):

In this part, we will experiment using other people's software. I suggest you use one of the SVM implementations available at http://www.support-vector-machines.org/SVM_soft.html. There are also MATLAB SVM implementations that you can use, including the toolboxes in BRMLToolBox (<http://web4.cs.ucl.ac.uk/staff/d.barber/pmwiki/pmwiki.php?n=Brml.Software>) and PMTK3 (<https://github.com/probml/pmtk3>) with discussions (<https://code.google.com/p/pmtk3/>). You might need to transform the data format.

Using the same binary data set, train the following SVMs (using just the training data): a linear SVM, and an RBF SVM (for **(extra credit 10pts)**). For the linear SVM, try different values of C ranging in 0.25, 0.5, 1, 2, 4. For the RBF SVM, try τ values (bandwidth) of 0.25, 0.5, 1, 2, 4. Plot error rates on both the development data and test data for the different values of C . How many support vectors are used for each model? Should this increase or decrease with C (why)?